

Diego Portfolio - Fullstack Architecture Document

Introduction

This document outlines the complete fullstack architecture for Diego Portfolio, including frontend implementation (React + Vite), backend services (FastAPI + AI), audio system design, and deployment infrastructure. It serves as the single source of truth for AI-driven development, ensuring consistency across the entire technology stack.

This unified approach combines frontend and backend architecture to streamline development for this modern fullstack application where these concerns are tightly integrated - particularly around the audio system, AI chatbot, and real-time visualizations.

Starter Template or Existing Project

Decision: No starter template - greenfield project built from scratch

Rationale:

- Unique requirements (audio visualizations, AI integration) don't align with standard templates
- Custom architecture needed for optimal performance
- Full control over dependencies and structure
- Educational value in building from foundation

Setup Approach:

```
bash

# Frontend
npm create vite@latest diego-portfolio -- --template react-ts

# Backend
mkdir backend && cd backend
python -m venv venv
source venv/bin/activate # On Windows: venv\Scripts\activate
pip install fastapi uvicorn langchain chromadb openai python-dotenv
```

Change Log

Date	Version	Description	Author
2025-01-XX	1.0	Initial architecture design	Winston (Architect)

High Level Architecture

Technical Summary

Diego Portfolio is a **serverless fullstack application** combining a static React SPA with serverless API functions. The frontend leverages Web Audio API for real-time audio analysis and Canvas/WebGL for visualizations, while the backend provides AI-powered conversational capabilities through FastAPI, LangChain, and a vector database.

The architecture prioritizes:

- **Performance:** Code splitting, lazy loading, optimized audio/visual rendering
- **Scalability:** Serverless functions scale automatically with traffic
- **Cost-effectiveness:** Static hosting + pay-per-request serverless ideal for portfolio traffic patterns
- **Developer Experience:** TypeScript across frontend, Python for AI/backend, clear separation of concerns

The system supports the PRD's five-track music player metaphor with seamless navigation, persistent audio playback, and an AI "DJ" that provides contextual information about Diego's background.

Platform and Infrastructure Choice

Selected Platform: Vercel (Frontend) + Vercel Serverless Functions (Backend API)

Rationale:

- **Unified deployment:** Single platform for frontend and backend simplifies CI/CD
- **Excellent DX:** Zero-config deployments, instant preview environments, automatic HTTPS
- **Optimal for React:** Built by Next.js team, excellent React/Vite support
- **Serverless Functions:** Native support for Python serverless functions
- **Edge Network:** Global CDN for fast asset delivery
- **Cost-effective:** Generous free tier, pay-as-you-grow pricing
- **TypeScript/Python native:** First-class support for both languages

Key Services:

- **Vercel Hosting:** Static site hosting with CDN
- **Vercel Serverless Functions:** Python API endpoints

- **Vercel Blob Storage:** Audio file storage (optional, or use external CDN)
- **External Vector DB:** Pinecone or hosted Chroma (Vercel functions connect to it)
- **Vercel Analytics:** Built-in performance monitoring

Deployment Regions:

- Primary: Global Edge (Vercel's CDN)
- Serverless: Auto-selected based on user location (Vercel Edge Network)

Alternative Consideration: If Vercel's free tier limits are exceeded, **Netlify** (similar offering) or **Firebase** (hosting + Cloud Functions) are suitable alternatives with minimal architecture changes.

Repository Structure

Structure: Monorepo

Monorepo Tool: pnpm workspaces (lightweight, fast, native to Vercel)

Package Organization:

```
diego-portfolio/
├── apps/
│   ├── web/          # React frontend
│   └── api/          # Python serverless functions
├── packages/
│   └── shared/       # Shared TypeScript types
├── docs/            # Project documentation
└── pnpm-workspace.yaml # Workspace configuration
└── package.json     # Root package.json
```

Rationale:

- **Monorepo benefits:** Shared types between frontend/backend, atomic commits, simplified CI/CD
- **pnpm advantages:** Efficient disk usage, fast installs, native Vercel support
- **Clear boundaries:** Separate apps for frontend/backend, shared package for common code
- **Easy local development:** Single repository, unified tooling

High Level Architecture Diagram

```
mermaid
```

```
graph TB
```

```
User[User Browser] -->|HTTPS| CDN[Vercel CDN / Edge Network]
```

```
CDN -->|Static Assets| Frontend[React SPA]
```

```
CDN -->|API Requests| API[Serverless Functions]
```

```
Frontend -->|Audio Files| AudioStorage[Audio Storage / CDN]
```

```
Frontend -->|Render| Canvas[Canvas/WebGL Visualizations]
```

```
Frontend -->|Play| AudioAPI[Web Audio API]
```

```
API -->|Query| VectorDB[Vector Database - Pinecone/Chroma]
```

```
API -->|Generate| LLM[OpenAI API]
```

```
VectorDB -->|Retrieve| Context[Diego's Knowledge Base]
```

```
Frontend -.Real-time.-> Visualizations[Audio Visualizations]
```

```
AudioAPI -.Analyze.-> Visualizations
```

```
style Frontend fill:#6366F1
```

```
style API fill:#10B981
```

```
style VectorDB fill:#F59E0B
```

```
style LLM fill:#EF4444
```

```
style Visualizations fill:#8B5CF6
```

Architectural Patterns

1. Jamstack Architecture: Static site generation with serverless APIs

- *Rationale:* Optimal performance, security, and scalability for content-heavy sites with dynamic API needs

2. Component-Based UI: Reusable React components with TypeScript

- *Rationale:* Maintainability, type safety, and clear component contracts across the application

3. Repository Pattern: Abstract data access for AI knowledge retrieval

- *Rationale:* Flexibility to swap vector database providers without changing business logic

4. API Gateway Pattern: Single FastAPI entry point for all backend operations

- *Rationale:* Centralized authentication, rate limiting, error handling, and logging

5. Real-time Data Processing: Web Audio API with requestAnimationFrame for visualizations

- *Rationale:* Smooth 60fps animations synchronized with audio playback

6. Retrieval-Augmented Generation (RAG): Vector search + LLM for contextual AI responses

- *Rationale:* Accurate, hallucination-resistant responses grounded in Diego's actual information
-

Tech Stack

This is the **DEFINITIVE** technology selection for the entire project. All development must use these exact versions.

Technology Stack Table

Category	Technology	Version	Purpose	Rationale
Frontend Language	TypeScript	5.3.x	Primary development language	Type safety, excellent tooling, prevents bugs
Frontend Framework	React	18.2.x	UI library	Component-based, mature ecosystem, team familiarity
Build Tool	Vite	5.0.x	Dev server and bundler	Fast HMR, optimal builds, native TypeScript support
UI Component Library	Headless UI	1.7.x	Accessible primitives	Unstyled, accessible, integrates with Tailwind
State Management	Zustand	4.4.x	Global state	Lightweight, simple API, TypeScript-first
Routing	React Router	6.21.x	Client-side routing	Industry standard, v6 improvements, loader pattern
Styling	TailwindCSS	3.4.x	Utility-first CSS	Rapid development, consistent design, JIT mode
HTTP Client	Axios	1.6.x	API requests	Interceptors, better error handling than fetch
Audio Processing	Web Audio API	Native	Audio analysis	Browser-native, powerful, no dependencies
Visualizations	Canvas API	Native	2D rendering	Performant, lightweight, sufficient for visualizations
Icons	Lucide React	0.303.x	Icon system	Consistent, tree-shakable, matches UX spec
Forms	React Hook Form	7.49.x	Form management	Minimal re-renders, validation, chat input

Category	Technology	Version	Purpose	Rationale
Backend Language	Python	3.11.x	Backend runtime	AI/ML ecosystem, FastAPI native, Diego's expertise
Backend Framework	FastAPI	0.108.x	API framework	Fast, async, automatic OpenAPI docs, type hints
AI Orchestration	LangChain	0.1.x	RAG pipeline	Modular AI components, vector store integrations
Vector Database	Pinecone	Latest	Embeddings storage	Managed, scalable, free tier, simple API
LLM Provider	OpenAI	1.6.x	Chat completion	GPT-4 Turbo for accuracy, embeddings for RAG
Validation	Pydantic	2.5.x	Data validation	Type-safe, FastAPI native, prevents errors
CORS	fastapi.middleware.cors	Built-in	Cross-origin requests	Enable frontend-backend communication
Environment	python-dotenv	1.0.x	Env variable loading	Secure config management
Frontend Testing	Vitest	1.1.x	Unit testing	Vite-native, fast, Jest-compatible API
Component Testing	React Testing Library	14.1.x	Component tests	Best practices, user-centric testing
Backend Testing	pytest	7.4.x	Python unit/integration tests	Standard Python testing, async support
E2E Testing	Playwright	1.40.x	End-to-end tests	Cross-browser, reliable, great DX
Linting	ESLint	8.56.x	Code quality (JS/TS)	Catch errors, enforce standards
Formatting	Prettier	3.1.x	Code formatting	Consistent style, automatic
Package Manager	pnpm	8.14.x	Dependency management	Fast, efficient, workspace support
Deployment	Vercel	Latest	Hosting platform	Frontend + serverless functions unified
Version Control	Git	Latest	Source control	Industry standard
CI/CD	GitHub Actions	Latest	Automation	Free for public repos, Vercel integration
Monitoring	Vercel Analytics	Built-in	Performance tracking	Zero-config, Core Web Vitals
Error Tracking	Sentry	7.91.x	Error monitoring	Frontend + backend, source maps

Data Models

Shared TypeScript Interfaces

These interfaces are defined in `(packages/shared/src/types/)` and shared between frontend and backend (FastAPI Pydantic models mirror these).

Track Model

typescript

```
// packages/shared/src/types/Track.ts

export type TrackId = 'university' | 'work' | 'projects' | 'skills' | 'hobbies';

export interface Track {
    id: TrackId;
    number: number;
    title: string;
    description: string;
    audioUrl: string;
    duration: number; // in seconds
    mood: string; // e.g., "Foundation & Learning", "Professional Growth"
    backgroundColor?: string; // Optional custom background color
}

export const TRACKS: Track[] = [
    {
        id: 'university',
        number: 1,
        title: 'University Years',
        description: 'Foundation & Learning',
        audioUrl: '/audio/track-1-university.mp3',
        duration: 180,
        mood: 'foundation',
    },
    // ... other tracks
];
```

Purpose: Represents each of the 5 portfolio tracks/sections

Relationships:

- Has associated content (rendered by track-specific components)

- Has audio file (played via AudioPlayer)
 - Referenced in navigation and playback state
-

Project Model

```
typescript

// packages/shared/src/types/Project.ts

export interface Project {
  id: string;
  title: string;
  description: string;
  longDescription?: string;
  techStack: string[]; // e.g., ['React', 'TypeScript', 'FastAPI']
  tags: string[]; // e.g., ['Web', 'AI', 'Full-Stack']
  imageUrl?: string;
  demoUrl?: string;
  githubUrl?: string;
  caseStudyUrl?: string;
  featured: boolean;
  completedDate: string; // ISO date string
}
```

Purpose: Represents Diego's projects showcased in Track 3

Relationships:

- Displayed as ProjectCards in Track 3
 - Clicked projects open ProjectDetailModal
 - Filterable by techStack or tags
-

ChatMessage Model

```
typescript
```

```

// packages/shared/src/types/Chat.ts

export type MessageRole = 'user' | 'assistant' | 'system';

export interface ChatMessage {
  id: string;
  role: MessageRole;
  content: string;
  timestamp: number; // Unix timestamp
  isExpanded?: boolean; // For two-tier responses
  detailedContent?: string; // Extended response content
  suggestedQuestions?: string[]; // Follow-up suggestions
}

export interface ChatConversation {
  id: string;
  messages: ChatMessage[];
  createdAt: number;
  updatedAt: number;
}

export interface ChatRequest {
  message: string;
  conversationHistory: ChatMessage[];
}

export interface ChatResponse {
  message: ChatMessage;
  suggestedQuestions: string[];
}

```

Purpose: Represents AI DJ chat messages and conversations

Relationships:

- ChatConversation contains array of ChatMessages
- ChatRequest sent to backend API
- ChatResponse received from backend, appended to conversation

AudioState Model

```
typescript
```

```
// packages/shared/src/types/Audio.ts

export type PlaybackState = 'playing' | 'paused' | 'loading' | 'error';

export interface AudioState {
  currentTrack: TrackId | null;
  playbackState: PlaybackState;
  currentTime: number; // seconds
  duration: number; // seconds
  volume: number; // 0-1
  isMuted: boolean;
  isShuffleOn: boolean;
  visualizationData: {
    frequencyData: Uint8Array;
    timeDomainData: Uint8Array;
  } | null;
}
```

Purpose: Global audio playback state managed by Zustand

Relationships:

- Consumed by AudioPlayer component
- Consumed by Visualization components
- Updated by audio event handlers

API Specification

REST API Endpoints

Base URL: <https://diego-portfolio.vercel.app/api>

POST /api/chat

Purpose: Send a message to the AI DJ chatbot

Request:

```
typescript
```

```
{  
  message: string;      // User's question  
  conversationHistory: ChatMessage[]; // Previous messages for context  
}
```

Response (200 OK):

```
typescript  
  
{  
  message: {  
    id: string;  
    role: 'assistant';  
    content: string;      // AI DJ's response (summary)  
    timestamp: number;  
    detailedContent?: string; // Expanded response (if applicable)  
    suggestedQuestions: string[];  
  };  
  suggestedQuestions: string[]; // Array of 3-5 follow-up questions  
}
```

Response (429 Too Many Requests):

```
typescript  
  
{  
  error: "Rate limit exceeded. Please wait before sending another message.";  
}
```

Response (500 Internal Server Error):

```
typescript  
  
{  
  error: "AI DJ is temporarily unavailable. Please try again later.";  
}
```

Authentication: None (public endpoint with rate limiting)

Rate Limiting: 10 requests per minute per IP address

Example Request:

```
bash
```

```
curl -X POST https://diego-portfolio.vercel.app/api/chat \  
-H "Content-Type: application/json" \  
-d '{  
  "message": "What programming languages does Diego know?",  
  "conversationHistory": []  
}'
```

GET /api/health

Purpose: Health check endpoint for monitoring

Response (200 OK):

```
typescript  
  
{  
  status: 'healthy';  
  timestamp: number;  
  version: string;  
}
```

Components

Frontend Component Architecture

The frontend is organized into a clear component hierarchy with separation of concerns.

Component: App

Responsibility: Root component, routing setup, global providers

Key Interfaces:

- Routes configuration
- Global state providers (Zustand, Theme)
- Error boundaries

Dependencies: React Router, Zustand stores

Technology Stack: React 18, TypeScript, React Router v6

Component: AudioPlayer

Responsibility: Global audio playback controls and state management

Key Interfaces:

- `play(): void` - Start playback
- `pause(): void` - Pause playback
- `seek(time: number): void` - Jump to specific time
- `setVolume(volume: number): void` - Adjust volume (0-1)
- `toggleMute(): void` - Mute/unmute
- `loadTrack(trackId: TrackId): void` - Load new track audio

Dependencies:

- Zustand audio store
- Web Audio API
- Browser HTMLAudioElement

Technology Stack: React, Web Audio API, Zustand

Implementation Note: Uses Web Audio API AudioContext for analysis while playing audio through HTMLAudioElement for reliability.

Component: Visualization

Responsibility: Real-time audio-reactive visual display using Canvas

Key Interfaces:

- Props: `frequencyData: Uint8Array`, `style: 'bars' | 'waveform' | 'radial'`
- Renders to Canvas element
- Updates at 60fps using requestAnimationFrame

Dependencies:

- Audio state from Zustand (frequency/time domain data)
- Canvas 2D rendering context

Technology Stack: React, Canvas API, requestAnimationFrame

Performance: Optimized rendering, only updates when audio playing, uses offscreen canvas if needed

Component: TrackView

Responsibility: Display individual track content with layout and navigation

Key Interfaces:

- Props: `trackId: TrackId`
- Renders track header, content, audio controls
- Handles track navigation

Dependencies:

- Track data (from TRACKS constant)
- AudioPlayer component
- Visualization component
- Navigation components

Technology Stack: React, React Router, TailwindCSS

Component: ChatInterface

Responsibility: AI DJ chat UI, message display, user input

Key Interfaces:

- `sendMessage(message: string): Promise<void>`
- Displays conversation history
- Handles message submission
- Shows typing indicator

Dependencies:

- Axios for API calls
- Chat state (local or Zustand)

Technology Stack: React, Axios, React Hook Form

Component: ProjectCard & ProjectDetailModal

Responsibility: Display project information, handle detail view

Key Interfaces:

- ProjectCard Props: `(project: Project)`, `(onClick: () => void)`
- Modal Props: `(project: Project)`, `(isOpen: boolean)`, `(onClose: () => void)`

Dependencies:

- Headless UI for modal
- Project data

Technology Stack: React, Headless UI, TailwindCSS

Backend Component Architecture

Component: ChatRouter (FastAPI Router)

Responsibility: Handle chat API endpoints

Key Interfaces:

- POST `/api/chat` endpoint
- Request validation (Pydantic models)
- Response formatting

Dependencies:

- ChatService
- Rate limiting middleware

Technology Stack: FastAPI, Pydantic

Component: ChatService

Responsibility: Business logic for AI chat interactions

Key Interfaces:

- `async process_message(message: str, history: List[Message]) -> ChatResponse`
- Orchestrates RAG pipeline
- Formats responses

Dependencies:

- VectorStoreService
- LLMService

Technology Stack: LangChain, Python async

Component: VectorStoreService

Responsibility: Vector database operations (embed, store, query)

Key Interfaces:

- `async query(text: str, k: int) -> List[Document]`
- Retrieves relevant context from knowledge base

Dependencies:

- Pinecone client
- OpenAI embeddings

Technology Stack: LangChain, Pinecone, OpenAI

Component: LLMService

Responsibility: LLM API calls and response generation

Key Interfaces:

- `async generate_response(query: str, context: str, history: List[Message]) -> str`
- Constructs prompts with system instructions

- Handles streaming responses (if needed)

Dependencies:

- OpenAI API client
- Prompt templates

Technology Stack: LangChain, OpenAI GPT-4 Turbo

Component Interaction Diagram

```

mermaid
sequenceDiagram
    participant User
    participant Frontend
    participant AudioPlayer
    participant Visualization
    participant ChatInterface
    participant API
    participant VectorDB
    participant LLM

    User->>Frontend: Navigate to Track 2
    Frontend->>AudioPlayer: loadTrack('work')
    AudioPlayer->>AudioPlayer: Load audio file
    AudioPlayer->>Visualization: Emit frequency data
    Visualization->>Visualization: Render canvas (60fps)

    User->>ChatInterface: Type question
    ChatInterface->>API: POST /api/chat
    API->>VectorDB: Query relevant docs
    VectorDB-->>API: Return context
    API->>LLM: Generate response with context
    LLM-->>API: Return AI message
    API-->>ChatInterface: ChatResponse
    ChatInterface->>ChatInterface: Display message
  
```

Database Schema

Vector Database (Pinecone)

Index Name: diego-portfolio-knowledge

Dimensions: 1536 (OpenAI text-embedding-3-small)

Metric: Cosine similarity

Document Structure:

```
python
{
  "id": "doc_123",      # Unique document ID
  "values": [0.1, 0.2, ...], # 1536-dim embedding vector
  "metadata": {
    "text": str,        # Original text chunk
    "source": str,      # e.g., "resume", "project_xyz"
    "category": str,    # e.g., "education", "work", "skills"
    "date": str,        # ISO date if applicable
  }
}
```

Knowledge Base Categories:

- **resume:** Career summary, work history
- **education:** University, degrees, coursework
- **projects:** Project descriptions, technologies, outcomes
- **skills:** Technical skills with proficiency levels
- **interests:** Hobbies, music, personal interests

Chunking Strategy:

- Max chunk size: 500 tokens
- Overlap: 50 tokens
- Preserve semantic boundaries (don't split mid-sentence)

Query Process:

1. User question embedded with OpenAI
2. Vector similarity search (top-k=5)
3. Retrieved chunks passed as context to LLM

4. LLM generates grounded response

Source Tree

Unified Project Structure

```
diego-portfolio/
├── .github/
│   └── workflows/
│       ├── ci.yml          # Run tests on PR
│       └── deploy.yml       # Deploy to Vercel on merge
|
└── apps/
    ├── web/                # Frontend React application
    │   └── public/
    │       ├── audio/         # Audio files for each track
    │       │   ├── track-1-university.mp3
    │       │   ├── track-2-work.mp3
    │       │   ├── track-3-projects.mp3
    │       │   ├── track-4-skills.mp3
    │       │   └── track-5-hobbies.mp3
    │       ├── images/         # Static images
    │       └── favicon.ico
    |
    └── src/
        ├── components/      # Reusable UI components
        │   └── audio/
        │       ├── AudioPlayer.tsx
        │       ├── AudioControls.tsx
        │       └── Visualization.tsx
        │
        └── chat/
            ├── ChatInterface.tsx
            ├── ChatMessage.tsx
            └── ChatInput.tsx
        └── navigation/
            ├── TrackList.tsx
            ├── TrackNavigation.tsx
            └── ShuffleButton.tsx
        └── projects/
            ├── ProjectCard.tsx
            ├── ProjectGrid.tsx
            └── ProjectDetailModal.tsx
```

```
|- |- |- tracks/
|- |- |- | TrackCard.tsx
|- |- |- | TrackHeader.tsx
|- |- ui/      # Shared UI primitives
|- |- | Button.tsx
|- |- | Modal.tsx
|- |- | Card.tsx

|- |- pages/      # Page components/routes
|- |- | AlbumView.tsx  # Landing page
|- |- | TrackView.tsx  # Individual track pages
|- |- | NotFound.tsx  # 404 page

|- |- hooks/      # Custom React hooks
|- |- | useAudioPlayer.ts
|- |- | useVisualization.ts
|- |- | useChat.ts
|- |- | useKeyboardShortcuts.ts

|- |- stores/      # Zustand state stores
|- |- | audioStore.ts  # Audio playback state
|- |- | chatStore.ts  # Chat conversation state

|- |- services/    # API client services
|- |- | chatService.ts  # Axios API calls

|- |- utils/       # Utility functions
|- |- | audioUtils.ts  # Audio processing helpers
|- |- | visualizationUtils.ts
|- |- | formatters.ts  # Date, time formatting

|- |- styles/      # Global styles
|- |- | globals.css  # Tailwind imports, custom CSS

|- |- data/        # Static data
|- |- | tracks.ts   # TRACKS constant
|- |- | projects.ts  # PROJECTS constant
|- |- | skills.ts    # SKILLS constant

|- |- App.tsx      # Root component
|- |- main.tsx     # Entry point
|- |- vite-env.d.ts # Vite type declarations

|- |- tests/       # Frontend tests
```

```
    |   |   |   └── unit/          # Component unit tests
    |   |   |   └── integration/   # Integration tests
    |   |   └── e2e/            # Playwright E2E tests
    |
    |   └── .env.example      # Environment variables template
    |   └── .env.local        # Local env (gitignored)
    |   └── index.html       # HTML entry point
    |   └── package.json
    |   └── tsconfig.json     # TypeScript config
    |   └── tsconfig.node.json
    |   └── vite.config.ts    # Vite configuration
    |   └── tailwind.config.js # Tailwind configuration
    |   └── postcss.config.js
    |   └── README.md
    |
    └── api/                  # Backend Python serverless functions
        ├── chat.py           # POST /api/chat endpoint
        └── health.py         # GET /api/health endpoint
    |
    └── services/
        ├── chat_service.py   # Chat business logic
        ├── vector_store_service.py # Pinecone operations
        └── llm_service.py     # OpenAI LLM calls
    |
    └── models/
        └── schemas.py        # Pydantic models
    |
    └── utils/
        ├── prompts.py        # System prompt templates
        └── rate_limit.py     # Rate limiting logic
    |
    └── scripts/
        ├── ingest_knowledge.py # Script to populate vector DB
        └── test_embeddings.py  # Test vector search
    |
    └── data/
        ├── resume.txt        # Source data for embeddings
        ├── projects.txt
        ├── skills.txt
        └── interests.txt
    |
    └── requirements.txt     # Python dependencies
    └── .env.example
    └── .env                  # Local env (gitignored)
```

```
|   └── README.md  
  
|  
|   └── packages/  
|       └── shared/          # Shared TypeScript types  
|           └── src/  
|               └── types/  
|                   ├── Track.ts  
|                   ├── Project.ts  
|                   ├── Chat.ts  
|                   └── Audio.ts  
|  
|           └── package.json  
|           └── tsconfig.json  
|           └── README.md  
  
|  
|   └── docs/                # Project documentation  
|       ├── prd.md            # Product Requirements Document  
|       ├── front-end-spec.md # UI/UX Specification  
|       ├── architecture.md   # This document  
|       └── brainstorming-session-results.md  
  
|  
|   └── scripts/             # Build/utility scripts  
|       └── setup-local.sh    # Local development setup  
  
|  
|   └── .gitignore  
|   └── .prettierrc          # Prettier configuration  
|   └── .eslintrc.json        # ESLint configuration  
|   └── pnpm-workspace.yaml  # pnpm workspace config  
|   └── package.json          # Root package.json  
|   └── vercel.json           # Vercel deployment config  
|   └── README.md             # Project overview
```

Development Workflow

Local Development Setup

Prerequisites

```
bash
```

```
# Required software
- Node.js 18+ (with npm/npx)
- pnpm 8+ (npm install -g pnpm)
- Python 3.11+
- Git
```

Initial Setup

```
bash

# Clone repository
git clone https://github.com/diego/portfolio.git
cd diego-portfolio

# Install frontend dependencies
pnpm install

# Set up backend
cd apps/api
python -m venv venv
source venv/bin/activate # Windows: venv\Scripts\activate
pip install -r requirements.txt

# Copy environment files
cp apps/web/.env.example apps/web/.env.local
cp apps/api/.env.example apps/api/.env

# Configure environment variables (see below)
# Edit .env.local and .env with your values
```

Environment Configuration

Frontend (.env.local):

```
bash

# API endpoint (local development)
VITE_API_URL=http://localhost:8000/api

# Analytics (optional for local)
VITE_VERCEL_ANALYTICS_ID=
```

Backend (.env):

```
bash

# OpenAI API
OPENAI_API_KEY=sk-...

# Pinecone
PINECONE_API_KEY=...
PINECONE_ENVIRONMENT=...
PINECONE_INDEX_NAME=diego-portfolio-knowledge

# Application
ENVIRONMENT=development
ALLOWED_ORIGINS=http://localhost:5173,http://localhost:3000

# Rate Limiting
RATE_LIMIT_PER_MINUTE=10
```

Development Commands

Start Frontend (from root):

```
bash

pnpm --filter web dev
# Runs on http://localhost:5173
```

Start Backend (from apps/api):

```
bash

source venv/bin/activate
uvicorn chat:app --reload --port 8000
# Runs on http://localhost:8000
```

Run All Tests:

```
bash
```

```
# Frontend tests  
pnpm --filter web test  
  
# Backend tests  
cd apps/api && pytest  
  
# E2E tests  
pnpm --filter web test:e2e
```

Lint and Format:

```
bash  
  
# Lint all code  
pnpm lint  
  
# Format all code  
pnpm format
```

Deployment Architecture

Deployment Strategy

Frontend Deployment:

- **Platform:** Vercel
- **Build Command:** (pnpm --filter web build)
- **Output Directory:** (apps/web/dist)
- **CDN/Edge:** Vercel Edge Network (global)

Backend Deployment:

- **Platform:** Vercel Serverless Functions
- **Runtime:** Python 3.11
- **Function Location:** (apps/api/*.py)
- **Deployment Method:** Automatic via Vercel CLI or Git integration

Deployment Triggers:

- Push to `main` branch → Production deployment
- Pull requests → Preview deployments
- Manual deploys via Vercel CLI

CI/CD Pipeline

GitHub Actions Workflow:

yaml

```
# .github/workflows/ci.yml
name: CI

on:
  pull_request:
    branches: [main, develop]
  push:
    branches: [main]

jobs:
  test-frontend:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v4
      - uses: pnpm/action-setup@v2
        with:
          version: 8
      - uses: actions/setup-node@v4
        with:
          node-version: 18
          cache: 'pnpm'
      - run: pnpm install
      - run: pnpm --filter web test
      - run: pnpm --filter web build

  test-backend:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v4
      - uses: actions/setup-python@v4
        with:
          python-version: '3.11'
      - run: cd apps/api && pip install -r requirements.txt
      - run: cd apps/api && pytest

  e2e:
    runs-on: ubuntu-latest
    needs: [test-frontend, test-backend]
    steps:
      - uses: actions/checkout@v4
      - uses: pnpm/action-setup@v2
      - uses: actions/setup-node@v4
```

- **run:** pnpm install
- **run:** pnpm --filter web test:e2e

Vercel Deployment:

```
json
// vercel.json
{
  "buildCommand": "pnpm install && pnpm --filter web build",
  "devCommand": "pnpm --filter web dev",
  "installCommand": "pnpm install",
  "framework": "vite",
  "outputDirectory": "apps/web/dist",
  "functions": {
    "apps/api/*.py": {
      "runtime": "python3.11"
    }
  },
  "rewrites": [
    {
      "source": "/api/:path*",
      "destination": "/apps/api/:path*"
    }
  ]
}
```

Environments

Environment	Frontend URL	Backend URL	Purpose
Development	http://localhost:5173	http://localhost:8000	Local development
Preview	https://diego-portfolio-*-.preview.vercel.app	Same (serverless)	PR previews, testing
Production	https://diego-portfolio.vercel.app	Same (serverless)	Live environment

Security and Performance

Security Requirements

Frontend Security:

- **Content Security Policy (CSP):** Configure in Vercel headers

```
Content-Security-Policy: default-src 'self'; script-src 'self' 'unsafe-inline'; style-src 'self' 'unsafe-inline'; img-src 'self' data:  
https://; media-src 'self' https://; connect-src 'self' https://diego-portfolio.vercel.app;
```

- **HTTPS Only:** Enforced by Vercel
- **XSS Prevention:** React's built-in escaping, avoid dangerouslySetInnerHTML
- **CORS Configuration:** Backend only accepts requests from production domain

Backend Security:

- **API Rate Limiting:** 10 requests/minute per IP (configurable)

```
python  
  
from fastapi import Request  
from slowapi import Limiter  
from slowapi.util import get_remote_address  
  
limiter = Limiter(key_func=get_remote_address)  
  
@app.post("/api/chat")  
@limiter.limit("10/minute")  
async def chat_endpoint(request: Request, ...):  
    ...
```

- **Input Validation:** Pydantic models validate all inputs
- **Secrets Management:** Environment variables, never in code
- **CORS Policy:** Whitelist production domains only

```
python  
  
app.add_middleware(  
    CORSMiddleware,  
    allow_origins=["https://diego-portfolio.vercel.app"],  
    allow_methods=["POST", "GET"],  
    allow_headers=["Content-Type"],  
)
```

Data Protection:

- **No Sensitive Data:** Portfolio is public, no user auth or PII

- **API Keys:** Stored as Vercel environment variables
- **Logs:** No sensitive information logged

Performance Optimization

Frontend Performance:

Bundle Size Target: < 200KB initial JS (gzipped)

Techniques:

- **Code Splitting:** Route-based splitting with React.lazy

typescript

```
const TrackView = lazy(() => import('./pages/TrackView'));
const ChatInterface = lazy(() => import('./components/chat/ChatInterface'));
```

- **Tree Shaking:** Vite automatically removes unused code
- **Asset Optimization:**
 - Images: WebP format, lazy loading, responsive srcset
 - Audio: MP3 compressed to 128-192kbps, progressive streaming
- **Caching:** Static assets cached with long TTL (Vercel CDN)

Runtime Performance:

- **Visualization Optimization:**

typescript

```
// Use RAF for smooth 60fps
const animate = () => {
  if (!isPlaying) return;

  analyser.getByteFrequencyData(frequencyData);
  drawVisualization(frequencyData);

  requestAnimationFrame(animate);
};
```

- **Memoization:** React.memo for expensive components
- **Virtual Scrolling:** Not needed (content fits on screen)

Backend Performance:

- **Response Time Target:** < 2 seconds for chat API
- **Vector Search:** Pinecone optimized for < 50ms queries
- **LLM Streaming:** Consider streaming responses for < 1s perceived latency
- **Caching:** Cache common questions (Redis if needed in future)

Monitoring:

- **Vercel Analytics:** Core Web Vitals tracked automatically
- **Sentry:** Error tracking for frontend + backend
- **Custom Metrics:**
 - API response times
 - Chat satisfaction (implicit via conversation length)
 - Track navigation patterns

Error Handling Strategy

General Approach

Error Model: Structured error responses with user-friendly messages

Exception Hierarchy:

- **Validation Errors:** 400 Bad Request (Pydantic handles this)
- **Rate Limit Errors:** 429 Too Many Requests
- **Server Errors:** 500 Internal Server Error (catch-all)

Error Propagation: Backend returns consistent error JSON, frontend displays user-friendly messages

Logging Standards

Library:

- Frontend: console (dev), Sentry (production)
- Backend: Python logging module + Sentry

Log Format:

```
python

# Backend
import logging

logging.basicConfig(
    level=logging.INFO,
    format='%(asctime)s - %(name)s - %(levelname)s - %(message)s'
)

logger = logging.getLogger(__name__)
```

Levels:

- **DEBUG:** Detailed technical info (only in development)
- **INFO:** General flow, successful operations
- **WARNING:** Potential issues, degraded performance
- **ERROR:** Errors that don't stop the service
- **CRITICAL:** Service-breaking errors

Required Context:

- Request ID (for tracing)
- User IP (for rate limiting debugging)
- Endpoint path
- Error stack traces (in production, sent to Sentry)

Error Handling Patterns

External API Errors (OpenAI, Pinecone)

Retry Policy:

```
python
```

```

from tenacity import retry, stop_after_attempt, wait_exponential

@retry(
    stop=stop_after_attempt(3),
    wait=wait_exponential(multiplier=1, min=2, max=10)
)
async def call_openai_api(prompt: str):
    # OpenAI API call
    ...

```

Circuit Breaker: Not needed initially (managed by Vercel timeouts)

Timeout Configuration: 30 seconds for LLM calls, 10 seconds for vector search

Error Translation:

```

python

try:
    response = await openai_client.chat.completions.create(...)
except openai.RateLimitError:
    raise HTTPException(429, "AI DJ is busy, please wait a moment")
except openai.APIError as e:
    logger.error(f'OpenAI API error: {e}')
    raise HTTPException(500, "AI DJ is temporarily unavailable")

```

Business Logic Errors

Custom Exceptions:

```

python

class VectorSearchError(Exception):
    """Raised when vector search fails"""
    pass

class LLMGenerationError(Exception):
    """Raised when LLM generation fails"""
    pass

```

User-Facing Errors:

- Friendly messages, no technical jargon

- Suggest actions ("Please try again" vs "Vector DB timeout")

Error Codes: HTTP status codes sufficient, no custom error codes

Data Consistency

Transaction Strategy: Not needed (stateless API, no database writes)

Idempotency: Chat requests are not idempotent (each generates new response)

Coding Standards

These standards are **MANDATORY** for AI agents. Focus on project-specific conventions and gotchas.

Core Standards

Languages & Runtimes:

- **Frontend:** TypeScript 5.3.x (strict mode enabled)
- **Backend:** Python 3.11.x (type hints required)

Style & Linting:

- **Frontend:** ESLint + Prettier (config in repo)
- **Backend:** Black formatter, flake8 linter

Test Organization:

- **Frontend:** `.test.tsx` files colocated with components
- **Backend:** `tests/` directory, mirror source structure

Critical Rules

Rule 1: No `localStorage` or `sessionStorage`

- **Why:** Not needed for this portfolio
- **Alternative:** Zustand persists to memory only during session

Rule 2: All API responses must use consistent error format

typescript

```
interface ApiError {  
  error: string;  
  details?: Record<string, any>;  
}
```

Rule 3: Audio visualizations must use requestAnimationFrame

- **Why:** Ensures smooth 60fps, doesn't block main thread
- **Never:** Use setInterval or setTimeout for animations

Rule 4: Never call LLM without context from vector DB

- **Why:** Prevents hallucinations, grounds responses in facts
- **Always:** Retrieve context first, then pass to LLM

Rule 5: All imports from shared package must use @shared alias

```
typescript  
// Correct  
import { Track } from '@shared/types/Track';  
  
// Incorrect  
import { Track } from '../../packages/shared/src/types/Track';
```

Rule 6: Component files must export default the main component

```
typescript  
// TrackView.tsx  
export default function TrackView() { ... }
```

Rule 7: Environment variables must be validated at startup

```
python
```

```

# Backend
from pydantic import BaseSettings

class Settings(BaseSettings):
    openai_api_key: str
    pinecone_api_key: str

    class Config:
        env_file = ".env"

settings = Settings() # Raises error if vars missing

```

Naming Conventions

Element	Frontend	Backend	Example
Components	PascalCase	-	AudioPlayer.tsx
Hooks	camelCase, prefix "use"	-	useAudioPlayer.ts
API Routes	-	snake_case	apps/api/chat.py
Functions	camelCase	snake_case	loadTrack(), process_message()
Constants	SCREAMING_SNAKE_CASE	SCREAMING_SNAKE_CASE	TRACKS, MAX_ATTEMPTS
Types/Interfaces	PascalCase	PascalCase	AudioState, ChatMessage

Test Strategy and Standards

Testing Philosophy

Approach: Test-after development (not strict TDD) with focus on critical paths

Coverage Goals: 70% overall, 90% for chat/audio logic

Test Pyramid:

- Unit tests: 70%
- Integration tests: 20%
- E2E tests: 10%

Test Types and Organization

Unit Tests

Framework: Vitest + React Testing Library (frontend), pytest (backend)

File Convention:

- Frontend: `ComponentName.test.tsx` colocated with component
- Backend: `tests/unit/test_service_name.py`

Location:

- Frontend: Colocated with source files
- Backend: `apps/api/tests/unit/`

Coverage Requirement: 70% minimum

AI Agent Requirements:

- Generate tests for all public functions
- Test happy path + error cases
- Mock all external dependencies (API calls, vector DB)
- Use AAA pattern (Arrange, Act, Assert)

Example Frontend Test:

```
typescript
```

```
// AudioPlayer.test.tsx
import { render, screen, fireEvent } from '@testing-library/react';
import AudioPlayer from './AudioPlayer';

describe('AudioPlayer', () => {
  it('plays audio when play button clicked', () => {
    render(<AudioPlayer />);
    const playButton = screen.getByRole('button', { name: /play/i });

    fireEvent.click(playButton);

    expect(screen.getByRole('button', { name: /pause/i })).toBeInTheDocument();
  });
});
```

Example Backend Test:

```
python
# tests/unit/test_chat_service.py
import pytest
from services.chat_service import ChatService

@pytest.mark.asyncio
async def test_process_message_returns_response():
    service = ChatService()

    response = await service.process_message("What are Diego's skills?", [])
    assert response.message.role == "assistant"
    assert len(response.message.content) > 0
    assert len(response.suggested_questions) >= 3
```

Integration Tests

Scope: Test interactions between components/services

Location:

- Frontend: [apps/web/tests/integration/](#)
- Backend: [apps/api/tests/integration/](#)

Test Infrastructure:

- **Frontend:** Mock Service Worker (MSW) for API mocking
- **Backend:** Test Pinecone index, mock OpenAI

Example:

```
typescript

// tests/integration/chat-flow.test.tsx
it('sends message to API and displays response', async () => {
  render(<ChatInterface />);

  const input = screen.getByPlaceholderText(/type your question/i);
  fireEvent.change(input, { target: { value: 'What projects?' } });
  fireEvent.click(screen.getByText(/send/i));

  await waitFor(() => {
    expect(screen.getByText(/Diego has built/i)).toBeInTheDocument();
  });
});
```

E2E Tests

Framework: Playwright

Scope: Critical user journeys

Location: `apps/web/tests/e2e/`

Test Data: Seed data or test environment

Key Scenarios:

1. Navigate between tracks
2. Play audio and see visualizations
3. Send chat message and receive response
4. Click project and view details
5. Use shuffle mode

Example:

```
typescript
```

```
// tests/e2e/navigation.spec.ts
import { test, expect } from '@playwright/test';

test('user can navigate between tracks', async ({ page }) => {
  await page.goto('http://localhost:5173');

  await page.click('text=Track 2: Work Experience');
  await expect(page).toHaveURL(/track\/work/);
  await expect(page.locator('h1')).toContainText('Work Experience');

  await page.click('text=Next');
  await expect(page).toHaveURL(/track\/projects/);
});

});
```

Test Data Management

Strategy:

- Unit tests: Mock data in test files
- Integration: Shared fixtures in `tests/fixtures/`
- E2E: Seeded test data or mock API

Fixtures:

```
typescript

// tests/fixtures/tracks.ts
export const mockTracks: Track[] = [
  {
    id: 'university',
    number: 1,
    title: 'University Years',
    // ... rest of mock data
  },
];
```

Continuous Testing

CI Integration: GitHub Actions runs all tests on PR

Performance Tests: Not critical for portfolio, monitor via Vercel Analytics

Security Tests: Dependency scanning via Dependabot

Next Steps

This architecture document is complete and ready for validation and development!

Immediate Actions

1. Review and Validate

- Product Owner validates architecture against PRD and UX spec
- Stakeholders approve technology choices and deployment plan

2. Set Up Development Environment

- Initialize repositories (frontend, backend, shared)
- Configure Vercel project
- Set up Pinecone vector database
- Obtain OpenAI API keys

3. Prepare AI Knowledge Base

- Collect Diego's resume, project descriptions, skills
- Write knowledge base documents (resume.txt, projects.txt, etc.)
- Run embedding script to populate Pinecone

4. Begin Story Creation

- Product Owner shards this document and PRD
- Scrum Master creates first stories from Epic 1
- Development begins with Foundation & Project Setup epic

Handoff Prompts

To Product Owner: "Architecture document complete. Please validate that:

- Technology stack aligns with PRD requirements
- All epics are technically feasible
- No critical architectural decisions are missing
- Deployment strategy matches project needs

Run the po-master-checklist to validate all documentation before development begins."

To Development Team: "Complete architecture available in docs/architecture.md. Key points:

- Monorepo with pnpm workspaces
- Frontend: React + Vite + TypeScript + Tailwind
- Backend: FastAPI + LangChain + Pinecone
- Deploy: Vercel (frontend + serverless functions)
- Audio: Web Audio API + Canvas visualizations
- AI: RAG pattern with vector search

Begin with Epic 1: Foundation & Project Setup. First story: Initialize Project with React + Vite."